

RevOnt: Reverse Engineering of Competency Questions from Knowledge Graphs

Fiorela Ciroku^{a,*}, Jacopo de Berardinis^b, Jongmo Kim^b, Albert Meroño-Peñuela^b,
Valentina Presutti^{a,*} and Elena Simperl^b

^bKing's College London, London, United Kingdom

^aAlma Mater Studiorum - University of Bologna, Bologna, Italy

ARTICLE INFO

Keywords:

knowledge engineering
knowledge graph
ontology development
competency question extraction

ABSTRACT

The process of developing ontologies - a formal, explicit specification of a shared conceptualisation - is addressed by well-known methodologies. As for any engineering development, its fundamental basis is the collection of requirements, which includes the elicitation of competency questions. Competency questions are defined through interacting with domain and application experts or by investigating existing datasets that may be used to populate the ontology i.e. its knowledge graph. The rise in popularity and accessibility of knowledge graphs provides an opportunity to support this phase with automatic tools. In this work, we explore the possibility of extracting competency questions from a knowledge graph. We describe in detail RevOnt, an approach that extracts and abstracts triples from a knowledge graph, generates questions based on triple verbalisations, and filters the questions to guarantee that they are competency questions. This approach is implemented utilising the Wikidata knowledge graph as a use case. The implementation results in a set of core competency questions from 20 domains present in the dataset presenting the knowledge graph, and their respective templates mapped to SPARQL query templates. We evaluate the extracted competency questions by calculating the BLEU score using human-annotated references. The results for the abstraction and question generation components of the approach show good to high quality. Meanwhile, the accuracy of the filtering component is above 86%, which is comparable to the state-of-the-art classifications.

1. Introduction

Knowledge engineering methodologies, especially for supporting ontology development, have been the subject of significant research in the Semantic Web field. Most methodologies (Peroni, 2017; Presutti, Daga, Gangemi and Blomqvist, 2009; Abdelghany, Darwish and Hefni, 2019; Abdelghany et al., 2019; John, Shah and Stewart, 2018; Schekotihin, Rodler, Schmid, Horridge and Tudorache, 2018; Paschke and Schäfermeier, 2018) include a requirement elicitation phase and an ontology evaluation/testing phase, as fundamental in their processes. Requirement elicitation is the process of extracting and collecting requirements from domain and application experts, stakeholders, datasets, etc., whereas the evaluation phase deals with assessing the quality of the ontology. The basis of both these phases is the requirements, in the form of Competency Questions (CQs). The competency questions play an important role in the ontology development life-cycle, as they represent the ontology requirements and are criteria for the evaluation of the ontology (Bezerra, Freitas and Santana, 2013). Despite

the key role of the CQs in the process, there is no systematic mechanism to elicit them. Moreover, even though there are several methodologies that place the competency questions at the centre of the ontology evaluation (Presutti et al., 2009; Bezerra et al., 2013), there is a notable need for tool support for this phase of the process (Fernández-Izquierdo and García-Castro, 2022).

Generally, CQs are defined in a top-down approach, but the potential coming from existing knowledge graphs (e.g. Wikidata (Vrandečić and Krötzsch, 2014), DBpedia (Auer, Bizer, Kobilarov, Lehmann, Cyganiak and Ives, 2007)) provides an opportunity to design a bottom-up approach to extract competency questions. The approach consists of reverse engineering an ontology, e.i. defining a method for extracting competency questions from knowledge graphs. This proposition is founded on the observation that recently, an increasing number of knowledge engineering projects (e.g. Polifonia¹, SPICE²) involve the construction of a knowledge graph, and in many cases, knowledge graphs are the starting point that leads to defining an ontology. A direct application of the method is being able to query the knowledge graph of origin in order to understand its content. Modern knowledge graphs are built using a combination of human and automated actions (e.g. information extraction from text, data input from databases). Usually, these actions are not well-documented and derive an ontology that is not adequately structured (Piscopo and Simperl, 2018). Therefore, understanding what competency questions a knowledge graph

*Corresponding author

✉ fiorela.ciroku2@unibo.it (F. Ciroku); Via San Giuseppe 8, 50122, Florence, Italy (F. Ciroku); valentina.presutti@unibo.it (V. Presutti)

🌐 <https://www.unibo.it/sitoweb/fiorela.ciroku2/en> (F. Ciroku);

<https://www.kcl.ac.uk/people/jacopo-de-berardinis> (J.d. Berardinis);

<https://www.kcl.ac.uk/people/albert-merono-penuela-1> (A.

Meroño-Peñuela); <https://www.unibo.it/sitoweb/valentina.presutti/en>

(V. Presutti); <https://www.kcl.ac.uk/people/elena-simperl> (E. Simperl)

ORCID(s): 0000-0002-3885-6280 (F. Ciroku); 0000-0001-6770-1969 (J.d. Berardinis); 0000-0002-4984-1674 (J. Kim); 0000-0003-4646-5842 (A.

Meroño-Peñuela); 0000-0002-9380-5160 (V. Presutti); 0000-0003-1722-947X

(E. Simperl)

¹<https://polifonia-project.eu/>

²<https://spice-h2020.eu/>

answers can support the development of its ontology, the knowledge graph's evaluation, documentation, and reuse.

Another application of the method can be to support the knowledge engineering process. We speculate that the requirement elicitation and the evaluation phases can be enhanced, both in terms of procedures and rules, and in terms of the quality of the outcomes, with such a method. Regarding the requirement elicitation, the data from a knowledge graph can provide a different view of the domain of the ontology that is being constructed, complementary to the view provided by domain experts that are involved in the process or other sources. This additional information can improve the coverage of the domain, and in turn, of the ontology model. At the same time, these competency questions, if properly formulated and documented, can later be used as criteria for the evaluation of the ontology.

Motivated by the need for tool support for the knowledge engineering process, the opportunities that knowledge graphs offer in terms of application, and by leveraging recent advances in natural language processing, we raise the following research questions:

- **RQ1:** Is it possible to extract competency questions directly from a knowledge graph?
- **RQ2:** What are the quality features that an extracted question must show to be classified as a competency question?

1.1. Background

The process for engineering an ontology starts with the requirement elicitation phase, as established in many methodologies such as Presutti et al. (2009); Abdelghany et al. (2019); John et al. (2018); Schekotihin et al. (2018); Paschke and Schäfermeier (2018). An initial set of requirements is formalized and prioritized by ontology engineers. Based on the priority given to the competency questions, ontology engineers develop an ontology iteratively. The competency questions are simultaneously used for the testing of the ontology to ensure its quality. While the ontology is built and tested, the engineers construct the respective knowledge graph. A synthesis of the main steps in an ontology development process is displayed in Figure 1.

Based on our experience with eXtreme Design (XD) Presutti et al. (2009), requirement elicitation is an engaging task that includes continuous interactions with experts and data investigation. In-person or online meetings with domain experts generally need substantial planning and time for the ontology engineer to thoroughly grasp the domain. Meanwhile, the current tool support, like Google forms and GitHub templates, is inadequate to provide a consistent framework for gathering information. The process still needs time and effort from the ontology developer to analyse the information, structure it, and comprehend the input. The same techniques and constraints apply to extracting competency questions from data sets. It is critical to emphasise that competency questions must be of a specific quality to be utilised for the intended purpose. A competency question's




Figure 1: A synthesis of the main steps in an ontology development process

fundamental quality check is that it must be translatable to a query that can validate the ontology. The transformation process from informal competency questions to formal competency questions lacks standardisation, is time-consuming, and requires a significant amount of effort.

An example of the requirement collection procedure based on eXtreme Design comes from the Polifonia project³. Firstly, we asked domain experts to create stories. In the context of the project, a story is a template for collecting requirements which might include information about the persona, the goal, the scenario, competency questions, and resources, described in detail by Bottini, Carriero, Carvalho, Cathé, Ciroku, Daga, Daquino, Davy-Rigaux, Guillotel-Nothmann, van Kemenade, Marzi, Meroño Peñuelala, Mulholland, Musumeci, Presutti and Scharnhorst (2021). Considering that the domain experts often are not knowledge engineers, the competency questions that we received required a considerable amount of manual work to be transformed into formal competency questions. For instance, an informal competency question was *"In which historical documents is there evidence of a musical composition?"*⁴. This competency question was followed up by additional interaction with the domain experts to understand what the concept of evidence means to them, which in turn was defined *"as any direct linguistic sign that refers to concepts (name, or part of it)"*. This information was used to reformulate the competency question and have a formal one as follows *"Which historical documents mention a musical composition?"*. In the testing phase, this competency question is translated into a SPARQL query based on the ontology that was developed. The manually drafted SPARQL query is shown in Example 1.1.

Example 1.1. A manually drafted SPARQL query

³The Polifonia project aims to support providers in revealing the content of value hidden in their digital objects, to support scholars in conducting research based on large, heterogeneous data sources and to enable citizens to access and understand Musical Heritage with all its complexity.

⁴The complete story is found in <https://github.com/polifonia-project/stories>

```

PREFIX mc: <...>
PREFIX rdf: <...>
SELECT DISTINCT ?historicalDocument
WHERE {
  ?composition a mc:MusicalComposition .
  ?composition mc:mentionedBy ?historicalDocument . }

```

This example demonstrates that the process that an ontology engineer has to complete for manually eliciting a competency question, formalising it, and transforming it to a query for evaluating the ontology is inefficient and demands substantial effort.

1.2. Our contribution

The main contributions of this article are twofold:

1. RevOnt, an approach for automatically extracting competency questions from a knowledge graph;
2. The RevOnt dataset, containing competency questions extracted from a subset of Wikidata.

RevOnt aims at reversing the ontology engineering process starting from the knowledge graph. The data from the knowledge graph is used to formulate questions with the means of Natural Language Processing (NLP) models. The questions are evaluated and filtered with techniques described in Section 3. The filtering results in a set of core competency questions that describe the part of the knowledge graph from where the triples were retrieved. These competency questions are used to generate respective SPARQL query templates, which in turn supports the testing of the ontology. An overview of the approach is shown in Figure 2.




Figure 2: An overview of the RevOnt approach

The knowledge graph that we use to experiment⁵ and evaluate the method is Wikidata, which is a collaboratively-edited multilingual knowledge graph. Wikidata focuses on objects that represent any topic, concept, or object. It is a popular knowledge graph, it reflects common knowledge on

⁵The implementation of the approach is found in <https://github.com/King-s-Knowledge-Graph-Lab/revont>.

many domains and provides resources that help to evaluate the approach such as ID, label, description, and statements. The limitations of this use case remain in the fact that the knowledge graph is not domain-specific. An example of a data model in Wikidata is shown in Figure 3⁶.




Figure 3: A datamodel in Wikidata. The datamodel provides key information for a Wikidata item such as identifier, label, description, aliases, statements, and references.

In the remainder of the paper, Section 2 summarises related works regarding ontology development methodologies with a focus on the concept of competency questions and related work on relevant fields of research such as ontology learning and schema induction. Section 3 describes the approach that RevOnt uses to extract competency questions from a knowledge graph. Section 4 presents the experimental setup and the results of the evaluation. Lastly, Sections 5 and 6 present the reflections of the work, provide indications for future work, and conclude the paper.

2. Related work

The early works in **ontology engineering** surfaced at the beginning of the '90s with projects such as TOVE by Fox, Barbuceanu and Gruninger (1996), Ontolingua by Gruber (1992), among others, described in (Jones, Bench-Capon and Visser, 1998). Building upon these foundational projects, later efforts, including the Enterprise Model Approach by Dietz (2006) and METHONTOLOGY by Corcho, Fernández-López, Gómez-Pérez and López-Cima (2005), emerged, drawing inspiration from and extending the pioneering works of the early '90s to the 2000s. These works, intending to define methodologies that are closer to engineering practices, are followed in the early 2000s by Sure,

⁶Picture retrieved from <https://www.wikidata.org/wiki/Q16222597>

Erdmann, Angele, Staab, Studer and Wenke (2002) and Öhgren and Sandkuhl (2005). After a decade, ontology development methodologies that include cyclic processes, continuous integration and testing as (Peroni, 2017; Presutti et al., 2009; Abdelghany et al., 2019; John et al., 2018; Schekotihin et al., 2018; Paschke and Schäfermeier, 2018) emerge and become a shared practice among engineers. One of the common denominators between these works is the concept of competency questions, early defined in Grüninger and Fox (1995). Based on several works such as (Lenat and Guha, 1993; Fadel, Fox and Gruninger, 1994; Uschold and King, 1995; Uschold and Gruninger, 1996; Presutti et al., 2009; Gangemi and Presutti, 2009; Vrandečić, 2009; Ren, Parvizi, Mellish, Pan, van Deemter and Stevens, 2014), the definition of competency question can be reasonably summarised as "a typical query that an expert might want to submit to a knowledge base of its target domain, for a certain task"⁷. The role of a competency question is twofold: (1) to drive the modelling of the ontology by serving as a requirement, (2) to assist the evaluation of the ontology by being expressed as a query. In our research, we use these criteria to evaluate the quality of the generated competency questions.

In consideration of the method and techniques that are used in our work, relevant related fields of research are ontology learning, and schema induction and discovery. **Ontology learning** is a multidisciplinary field that extracts terms, concepts, properties, and relationships from unstructured text using approaches from several disciplines such as knowledge representation, natural language processing, machine learning, etc. Surveys in ontology learning such as (Asim, Wasim, Khan, Mahmood and Abbasi, 2018), classify ontology learning techniques into three classes namely linguistic, statistical and logical. *Linguistic techniques* are based on language features and are commonly used for data preparation (speech tagging, parsing and lemmatization) as well as various other ontology learning tasks such as knowledge extraction. Prime examples of such techniques are Text2Onto⁸, CRCTOL⁹. *Statistical techniques* (C/NC value, contrastive analysis, clustering, co-occurrence analysis, term subsumption and ARM) rely entirely on statistics from the underlying corpus and overlook the underlying semantics. The majority of statistical approaches make substantial use of probabilities and are commonly utilised in early stages of ontology learning following linguistics preprocessing. Relevant tools that use statistical techniques are OntoGain (Frantzi, Ananiadou and Mima, 2000), OntoLearn (Kouagou, Heindorf, Demir and Ngonga Ngomo, 2022), ASIUM (Faure and Poibeau, 2000). Lastly, *Inductive logic programming* is a machine learning discipline that employs logic programming to generate hypotheses based on prior knowledge and a set of examples. Significant tools are Syndikate (Hahn and Romacker, 2001) and TextStorm¹⁰.

With the rapid growth of RDF and Knowledge Graph data resources, numerous studies on ontology learning using RDF have been proposed. These studies can be categorized based on the approach adopted for utilizing RDF in ontology learning. Some studies have leveraged RDF datasets to enhance the quality of ontology learning, given that ontologies may have limited instances and data values. Consequently, data-driven ontology learning techniques could easily encounter information scarcity issues. RDF datasets serve as a valuable resource for ontology learning due to their semi-structured nature and similarity to ontologies. Andrea and et al. proposed a possibilistic approach for testing OWL (Web Ontology Language) axioms against RDF facts, specifically for testing the 'SubClassOf' relationship (Tettamanzi, Faron-Zucker and Gandon, 2014). The ADOL was designed for automatic domain ontology learning from textual data with RDFs in general-purpose Knowledge Graphs. It utilizes RDF data to create semantic similarity with the ontology by comparing extracted words and relations to the RDF set (Chen and Gu, 2021). Conversely, Mid-Ontology has been introduced to automatically construct a simplified ontology from the RDF set of Linked Open Cloud (LOD) for the integration of different ontology schema (Zhao and Ichise, 2012).

Lastly, surveys such as (Ji, Qi, Gao and Wu, 2019) and (Kellou-Menouer, Kardoulakis, Troullinou, Kedad, Plexousakis and Kondylakis, 2022) present an overview of the state-of-the-art in **schema induction** and discovery, a research field dealing with the extraction or discovery of semantic schemas from unstructured or semi-structured data (Gick and Holyoak, 1983). The research is motivated by the fact that the data in the semantic web, either expressed in RDF or JSON, are not based on a predefined schema. The most widely used techniques in the schema discovery approaches are machine learning (classification, clustering and frequent pattern mining) and formal techniques (Formal Concept Analysis, bisimulation). The surveys discuss valuable works regarding implicit and explicit schema discovery approaches by taking into consideration the target problem, techniques, features, input, output, and quality aspects.

Relevant surveys such as (Pouriyeh, Allahyari, Kochut and Arabnia, 2018) and (Čebirić, Goasdoué, Kondylakis, Kotzinos, Manolescu, Troullinou and Zneika, 2019) describe **ontology summarisation** approaches that use centrality metrics (e.g. PageRank) to identify the most informative concepts/nodes or extract important subgraphs to facilitate query-testing for verifying requirements against accessible data. In contrast, a recent research related to the extraction of Common Conceptual Components¹¹ from multiple Ontologies using Ontology Design Patterns¹² is presented in (Asprino, Carriero and Presutti, 2021). The authors present a method that employs a non-extractive method to assist in the comprehension and comparison of

⁷Formulated by Gangemi and Presutti (2009)

⁸<http://neon-toolkit.org/wiki/1.x/Text2Onto.html>

⁹<http://nlp.cs.berkeley.edu/>

¹⁰<https://dwijottam-dutta.github.io/TextStorm/about/about.html>

¹¹A conceptual component (CC) is a complex (cognitive) relational structure that a designer implements in an ontology by using classes, properties, axioms, etc.

¹²http://ontologydesignpatterns.org/wiki/Main_Page

different ontologies. Starting with a corpus of ontologies, it uses community detection, word sense disambiguation, frame recognition, and clustering to automatically produce a catalogue of conceptual components and observable ontology design patterns. Further, in (Nuzzolese, Gangemi, Presutti and Ciancarini, 2011), the authors present a study for discovering Encyclopedic Knowledge Patterns (EKP)¹³ from Wikipedia¹⁴ page links. The patterns, according to the authors, may be used as lenses for exploring DBpedia¹⁵ or for developing new ontologies that inherit the data and textual grounding offered by DBpedia and Wikipedia. Data linking can also benefit from EKPs by modularizing the datasets to be linked.

Our research contributes to the field of knowledge engineering by enhancing the requirement elicitation and ontology testing processes. The competency questions extracted with our approach can be used to develop and test an ontology. Furthermore, extracting competency questions from a knowledge graph can indirectly aid ontology learning tasks, as this approach retrieves terms and relations through the use of natural language processing models and machine learning. For instance, a wealth of competency questions could assist in identifying the optimal and most suitable ontology by testing and validating the generated ontology throughout the ontology learning process. Lastly, RevOnt can support the extraction of a schema of a knowledge graph through its abstraction stage, where the natural language verbalisation of a triple is abstracted from the instance level to the class level.

3. The RevOnt approach

This section presents RevOnt, a framework for extracting competency questions from knowledge graphs. As depicted in Figure 4, this is divided into three stages: 1) verbalisation abstraction, 2) question generation, and 3) question filtering.

Verbalising a Knowledge Graph consists in generating grammatically correct natural language starting from interconnected triple-based claims – formed of subject, predicate, and object, Amaral (2022). Here, the purpose of the *Verbalisation Abstraction* stage is to transform the verbalisation from the instance level to the class level. For example, given the triple verbalisation "*Michael Jackson is a member of the Michael Jackson discography.*", without the abstraction the generated questions are: "*Who is a member of the Michael Jackson discography?*" and "*What is Michael Jackson a member of?*". These questions are not competency questions because they ask for specific instances, and not classes or properties. Thus, there is a need to abstract the verbalisation into a more general form. To complete the task, the *Verbalisation Abstraction* stage begins with a dataset, which contains the verbalisations of triples from a knowledge graph, as an input. The dataset should include data about the subject,

Table 1

A list of the language models, modules, datasets and services used in the RevOnt framework

Model	Category	Stage
WDV	Dataset	Input
MiniLM	LM	Verbalisation abstraction
Wikidata query service	Service	Verbalisation abstraction
T5	LM	Question generation
T5 Grammar Correction	LM	Question generation
BigCQ	Dataset	Question filtering
SBERT	LM	Question filtering
UMAP	Dim. reduction	Question filtering
HDBSCAN	Clustering	Question filtering

predicate, and object of the triple, descriptions of the instances, IDs that align them to a knowledge graph; and most importantly, the verbalisation of the triple. This information is necessary for the selection of the class in which the subject and the object of the triple is an instance of or a subclass of.

The second stage of the framework is *Question generation*. The aim of this stage is to generate three questions for each triple verbalisation. This choice of design is made to understand the types of competency questions that are generated when the method asks the model about different parts of the verbalisation sentence. Essentially, this stage generates questions based on the abstraction of the triple verbalisation (e.g. from "A music artist is a member of their own discography" to "Who is a member of a discography?"). Part of this stage is a grammar check task that corrects errors that are found in the questions.

Lastly, the *Question filtering* stage deals with the quality check and reduction of the questions generated in the previous step. The two main tasks of this stage are 1) mapping question templates to query templates, and 2) reduction of questions based on their similarity. The task of the mapping of question templates to query templates addresses the question "*When is a question a competency question?*".

We should look up papers by Jeff Pan 2014 and so-called competency question templates that will probably answer the question. Also look at the work by Valentina Tamma.

As discussed in Section 2, the answer is that a competency question must be used to query the ontology. Therefore, such a mapping would confirm whether the generated question is indeed a competency question. While the question reduction task filters redundant questions by searching for semantically equivalent competency questions. This stage is needed to identify a set of core competency questions from the whole set of questions that are generated in the previous stage.

Table 1 shows a list of NLP models, modules, datasets and services used for the implementation of the framework, their category and the stage where they are used.

¹³EKPs are Knowledge Patterns that are grounded in encyclopedic knowledge expressed as linked data and as natural language text.

¹⁴https://en.wikipedia.org/wiki/Main_Page

¹⁵<https://www.dbpedia.org/>




Figure 4: An overview of the RevOnt framework. The first stage, Verbalisation Abstraction, generates the abstraction of a triple verbalisation. The abstraction is used as input in the second stage, Question Generation, to generate three questions per triple and perform a grammar check. Lastly, the third stage, Question Filtering, filters the questions by performing different techniques.

3.1. The WDV dataset

There are several datasets that provide verbalisations of data such as the T-REx¹⁶ and WDV dataset¹⁷ for Wikidata entries, WebNLG¹⁸ for DBpedia¹⁹ entries, NYT-FB by Mintz, Bills, Snow and Jurafsky (2009) and FB15K-237 by Toutanova and Chen (2015) for Freebase²⁰, and so on.

A first implementation of the RevOnt framework was achieved for Wikidata, by leveraging the WDV dataset (Amaral, 2022). This dataset provides verbalisations of Wikidata claims and it contains 7.6K unique triples. According to Amaral, Rodrigues and Simperl (2022), WDV has considerably more entity types and predicates than comparable datasets, and it is intended to serve as a benchmark dataset for data verbalisation models used on Wikidata. WDV enables a tight coupling between single claims and text by directly connecting a triple-based claim to a natural language phrase. The WDV dataset contains verbalisation of claims from 20 different themes (or domains) from the most common like artist, sport teams, university to celestial body, chemical compounds, taxon, et cetera. The diversity of triple verbalisations in the dataset contributes to interesting results and imposes challenges as well. An example of a claim verbalisation from the *Airpot* theme is shown in Figure 5. The described claim verbalisation is structured in the form of key-value pairs. It contains RDF triple-related information including "*subject_label*" and "*subject_desc*", as well as claim-related information including "*theme_label*" and "*verbalisation*".

¹⁶<https://hadyelsahar.github.io/t-rex/>

¹⁷<https://github.com/gabrielmaia7/WDV>

¹⁸<https://gitlab.com/shimorina/webnlg-dataset>

¹⁹<https://www.dbpedia.org/>

²⁰<https://developers.google.com/freebase>

3.2. Verbalisation abstraction

The first stage of the RevOnt framework is the *Verbalization Abstraction*. Its role is to generate an abstraction of a triple verbalisation. The triple verbalisation of Knowledge Graphs refers to converting RDF triples into natural language text by utilizing their components and relationships. To perform the abstraction, it is necessary to recognize and categorize the entities present in the verbalisation. The initial intuition for this task was to use a Named Entity Recognition (NER) model. We experimented with language models such as Camembert-ner²¹, Camembert-base-multilingual-cased-ner-hrl²², Ner-english-large²³, Bio-Ner²⁴ and the spaCy library. The performance of each of these models was not satisfactory for 50% of the themes of the dataset.

Considering the performance of those language models on the WDV dataset, we designed a novel method to abstract the subject and object of a triple to the most similar-to-context Wikidata class. The procedure is shown in Algorithm 1.

For each triple, the method extracts the sentence embeddings of the descriptions of the subject and the object using the MiniLM language model²⁵. MiniLM is a sentence-level transformer model (Reimers and Gurevych, 2019) that maps sentences to a 384-dimensional dense vector space. By leveraging a pre-trained language model²⁶, MiniLM is fine-tuned on a number of datasets²⁷ using a contrastive

²¹<https://huggingface.co/Jean-Baptiste/camembert-ner>

²²<https://huggingface.co/Davlan/bert-base-multilingual-cased-ner-hrl>

²³<https://huggingface.co/flair/ner-english-large>

²⁴<https://github.com/library/bio-ner>

²⁵<https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2>

²⁶<https://huggingface.co/nreimers/MiniLM-L6-H384-uncased>

²⁷https://www.sbert.net/docs/pretrained_models.html

A claim verbalisation from the WDV dataset

```
{
  "claim_id": "Q20451714$49CB1F7A-FC51-41CF-81EA-3EC1EC81AB5D",
  "rank": "normal",
  "subject_id": "Q20451714",
  "property_id": "P366",
  "subject_label": "Livramento Do Brumado Airport",
  "property_label": "use",
  "object_label": "general aviation",
  "subject_desc": "airport in Brazil",
  "property_desc": "main use of the subject (includes current and former usage)",
  "object_desc": "civil use of aircraft excluding commercial transportation",
  "subject_alias": ["SNLB"],
  "property_alias": [{"function": "role", "mission": "purpose", "utility": "used for", "used in": "usage"}],
  "object_alias": [{"GA": "G/A", "genav": "civil aviation"}],
  "object_datatype": "wikibase-item",
  "object": {
    "value": {
      "entity-type": "item",
      "numeric-id": 1571929,
      "id": "Q1571929",
      "type": "wikibase-entityid"
    },
    "type": "wikibase-entityid"
  },
  "theme_root_class_id": "Q1248784",
  "theme_label": "Airport",
  "verbalisation": "Livramento Do Brumado Airport is used by general aviation.",
  "verbalisation_unk_replaced": "Livramento Do Brumado Airport is used by general aviation.",
  "sampling_weight": 2.642857143,
  "annotations": {
    "fluency_scores": [5, 5, 2, 4, 3],
    "fluency_mean": 3.8,
    "fluency_median": 4.0,
    "adequacy_scores": [1, 0, 0, 0, 0],
    "adequacy_majority_voted": 0.0,
    "adequacy_percentage": 0.0
  }
}
```

A depiction of triple-related key & value from a claim




Figure 5: A claim verbalisation from the WDV dataset. For the claim, the dataset provides the ID, rank, theme, and verbalisation of the claim. There are present also the label, description, aliases and ID of the subject, property and object of the triple.

Algorithm 1 Abstraction of a Wikidata triple

```
1: procedure TRIPLE VERBALISATION ABSTRACTION
2:   Create the sentence embedding of the subject description
3:   Create the sentence embedding of the object description
4:   Retrieve the Wikidata classes of the subject and object of the triple
5:   for each class do
6:     Get the synsets
7:     for each synset do
8:       Get the synset definition
9:       Create the sentence embedding of the synset definition
10:    Calculate the cosine similarity between the synset definition embedding and the subject/object descriptions' embeddings
11:   return Most similar synset
```

objective. Intuitively, a contrastive loss is used to minimise the cosine similarity between similar sentence pairs, while maximising that of other sentences in the same batch.

In the next step, RevOnt selects the English label of Wikidata classes where an entity is an instance of or a subclass of. The query that we use for selecting this information from the knowledge graph is shown below. For this

task, we use the Wikidata Query Service²⁸. Wikidata Query Service is an implementation of a SPARQL server, based on Blazegraph²⁹ engine, to service queries for Wikidata and other data sets.

```
SELECT DISTINCT ?cLabel
WHERE {{
  wd:{id} wdt:P31/wdt:P279? ?c .
  ?c rdfs:label ?cLabel .
}}
FILTER(LANG(?cLabel) = "en") }}
```

For instance, for the triple verbalisation "*Michael Jackson is a member of the Michael Jackson discography.*", according to the dataset shown in Figure 5, the subject is *Michael Jackson* and the object is *Michael Jackson discography*. The respective descriptions are *American recording artist; singer and songwriter (1958-2009)* and *Wikimedia artist discography*. RevOnt extracts the sentence embeddings of the subject and object description and queries the Wikidata knowledge graph to select the classes of the subject and object of the triple. The result for the subject is shown in Example 3.1.

Example 3.1. The Wikidata classes for the subject "Michael Jackson"

```
['human', 'natural person', 'omnivore', 'person',
'mammal', 'Homo sapiens']
```

Once RevOnt retrieves the classes, for each class it gets the corresponding Wordnet synsets. Wordnet is a large electronic lexical database for English (Fellbaum, 2010). Cognitive synonyms (synsets) are groups of nouns, verbs, adjectives, and adverbs that each communicate a separate notion. Synsets are linked together via conceptual-semantic and linguistic relationships. In Example 3.2 we present the respective synsets of the classes where "*Michael Jackson* is an instance of/subclass of. For each synset, it retrieves the definition and computes its sentence embeddings by using the MiniLM model.

Example 3.2. The synsets of the Wikidata classes

```
{[Synset('homo.n.02')],
[],
[Synset('omnivore.n.01'), Synset('omnivore.n.02')],
[Synset('person.n.01'), Synset('person.n.02'),
Synset('person.n.03')], [Synset('mammal.n.01')],
[]}
```

In the last step, RevOnt calculates the cosine similarity between the embeddings of the definition of the synsets of each class with the embeddings of the description of the subject or object. The cosine similarity is calculated with the help of the Natural Language Toolkit (NLTK). NLTK is a suite of Python modules providing many NLP data types, processing tasks, corpus samples, and readers, together with

²⁸<https://query.wikidata.org/>

²⁹https://github.com/blazegraph/database/wiki/Main_Page

animated algorithms, tutorials, and problem sets Loper and Bird (2002). For the example shown above, the Wikidata class that is the most similar to the description of the subject *Michael Jackson is human*. As for the object of the triple, the Wikidata class is *discography*.

The values that the algorithm returns populate a Python dictionary that is used for the abstraction task. More specifically, the subjects/objects and their corresponding most similar-to-description class are added to a dictionary, as shown in Example 3.3.

Example 3.3. *The pattern dictionary*

```
{'Michael Jackson': 'human',  
'Michael Jackson Discography': 'discography'}
```

This dictionary is used to replace the subject and the object in a triple with the respective class. To continue on the same example above, the result of the verbalisation abstraction is presented in Example 3.4. This concludes the first stage of the framework, and the abstracted verbalisations are passed on to the second stage, Question generation.

Example 3.4. *The abstraction of the verbalisation*

Verbalisation: Michael Jackson is a member of the
Michael Jackson discography.
Abstraction: Human is a member of the discography.

3.3. Question generation

After the abstraction of the verbalisation, the next stage of the approach is to generate questions. Given a triple verbalisation where entities have been generalised, the goal is to automatically generate a set of questions addressing the content of the verbalisation. For this task, we have used the Text-to-Text Transfer Transformer (T5) (Roberts and Raffel, 2020) fine-tuned on the Stanford Question Answering Dataset (SQuAD) (Rajpurkar, Zhang, Lopyrev and Liang, 2016) for question generation. This was achieved by prepending the answer to the context. The specific T5 instance we used is labelled as `t5-base-finetuned-question-generation-ap`³⁰. Notably, Romero (2021) demonstrated that this model can be handle generative tasks such as abstractive summarisation, classification tasks such as natural language inference, and even regression tasks.

To generate questions, the model requires as input a context (sentence) and an answer. When an answer is not provided, the model will generate a question that is answered by the object of the sentence. For each triple, we have provided the model with the abstraction of the verbalisation as the context, and three answers: the class of subject, the property, and the class of the object. In most observed cases, when the class of subject and the class of the object are the answers, the questions that are generated are inverse in relation to each other; hence they may provide complementary requirements for the triple. As for the cases when the property is the answer, the questions that are generated are

often the same as the ones when the answer is the class of the object, but there are also many cases when the questions are quite interesting and not so straightforward. This behaviour of the model can be explained with the distance between the triple and the verbalisation. As noticed in the example below, while the property of the triple is *discography*, this property is not present in the verbalisation. This is the case for the majority of the verbalisations present in the dataset. In Example 3.5, we show the results of the Question Generation stage given the abstraction from Example 3.4.

Example 3.5. *Question generation*

Context: Human is a member of the discography.
Answer 1: human
Question 1: What is a member of the discography?
Answer 2: discography
Question 2: What is the relationship between
a human and a discography?
Answer 3: discography
Question 3: What is a human a member of?

Once the questions are generated, RevOnt performs a grammar check to detect and fix errors. This task is performed by the T5 Grammar Correction model³¹. Trivially, the model generates a revised version of the given text with the goal of addressing grammatical errors. It is trained with Happy Transformer³² using a the JFLEG dataset (Napoles, Sakaguchi and Tetreault, 2017).

3.4. Question Filtering

In this section we cover the third and last stage of RevOnt, which consists in the filtering of the questions that are generated from the second stage. This stage is divided into two parts: (i) *matching* competency questions to SPARQL templates (to address quality issues of CQs); and (ii) *question reduction* through paraphrase detection (to remove equivalent questions) and question clustering (to identify groups of similar requirements).

3.4.1. CQ-query template mapping

To address **RQ2**, regarding the quality of the extracted questions from a knowledge graph, we filter the questions that are generated from the abstraction of the verbalisation with the help of the BigCQ dataset³³ by Wiśniewski, Potoniec and Ławrynowicz (2021). The dataset contains a set of Competency Question templates paired with SPARQL-OWL query templates, and it has been built upon previous work, CQ2SPARQLOWL by Wiśniewski, Potoniec, Ławrynowicz and Keet (2018). The method and dataset in CQ2SPARQLOWL are a benchmark of schema-level Competency Questions and corresponding SPARQL-OWL patterns. Therefore, in this task we select the questions that are mapped to SPARQL queries, proving that they are indeed competency questions. To perform this task, we need to

³¹<https://huggingface.co/vennify/t5-base-grammar-correction>

³²<https://github.com/EricFillion/happy-transformer>

³³<https://github.com/dwisniewski/BigCQ>

³⁰<https://huggingface.co/mrm8488/t5-base-finetuned-question-generation-ap>

transform questions into templates. The transformation is shown in Example 3.6.

Example 3.6. *The questions to question template transformation*

Question-1: What is a member of the discography?

Question template-1: What is EC1 of EC2?

Question-2: How many sub-genres of rock have the same defining characteristics as rock itself?

Question template-2: How many specializations of something being equal to it exist?

The BigCQ dataset comprises multiple documents, each containing a variety of question templates corresponding to a SPARQL-OWL pattern. There is a notable difference between the generated CQs and the SPARQL-OWL patterns, which makes a direct comparison challenging. This is because CQs are expressed in natural language, whereas SPARQL-OWL patterns are formulated using a SQL-based query syntax. The question templates serve as a useful intermediary between the CQs and SPARQL-OWL patterns. They possess a structure similarity to the CQs and have a mapping relation to SPARQL-OWL patterns. The procedure for mapping CQs to query templates within the BigCQ dataset unfolds as follows: Initially, both generated CQs and question templates are converted into embedding vectors utilizing the sentence transformer. Secondly, the average L2-norm distance between the embedding of a generated CQ and the set of embeddings of question templates from the same document within the BigCQ dataset is computed, as expressed in the equation below.

$$Dist_2(cq_i, D_k) = \frac{1}{|\{qt_j \in D_k\}|} \sum_{j, qt_j \in D_k} \|cq_i - qt_j\|_2^2 \quad (1)$$

where cq_i is i -th generated CQ, qt_j is j -th query template, and D_k is k -th document containing multiple query templates and mapping SPARQL-OWL pattern of the BigCQ dataset.

Thirdly, the document representing the minimum distances with a generated CQ and its corresponding SPARQL-OWL patterns ($\text{argmin}_k Dist_2(cq_i, D_k)$) is selected. This step is reiterated for all generated CQs and documents within the BigCQ dataset, yielding a mapping list of CQs and SPARQL-OWL patterns. The next step is to add the new templates (derived from generated CQs) to the BigCQ dataset and regenerate the dataset. The result of the regeneration is a file with a list of unique competency question templates, a file with a list of unique SPARQL-OWL query templates, and a folder with mappings from SPARQL-OWL templates to CQ templates. Each file represents a separate SPARQL-OWL query template in a JSON document following the following schema:

Example 3.7. *Question template to SPARQL query template mapping*

```
{ 'query': 'ASK WHERE <c1> rdfs:subClassOf
[ rdf:type owl:Restriction ; owl:onProperty <op1> ;
```

```
owl:someValuesFrom [ rdf:type owl:Restriction ;
owl:onProperty <op2> ; owl:hasValue <c2> ] ] . ',
'cqs': ['is there any c1 that can be counted as a sort of
something that op1 something that op2 EC2?', 'Is it true
that EC1 is something that op1 something that op2 EC2?',
...] }
```

From the new dataset, we are able to select the competency questions which have templates mapped to SPARQL queries.

3.4.2. Question reduction

As some of the questions generated in the previous steps may be redundant, or show negligible semantic variations that are of little interest to ontology engineers, we filter out questions based on their similarity. This has two potential benefits: (i) it mitigates the noise and the artifacts introduced in the previous steps (e.g. questions with unclear or inconsistent semantics); (ii) it reduces the number of competency questions that will be presented to the ontology engineers. Depending on the desired level of filtering, here we provide two methods to identify groups of related questions.

- *Similarity grouping*, which aims at detecting groups (or clusters) of semantically similar – yet not necessarily identical, questions.
- *Paraphrase detection*, a specialisation of the former task focused on detecting questions with the same exact meaning (e.g. “Is Batman a friend of Robin?” and “Is Batman Robin’s friend?”).

If the goal is to retain the largest number of questions, for example, because the ontology should model the domain at a granular level, the latter filtering method is more indicated. Instead, similarity grouping operates a more drastic reduction based on the relatedness of questions; which is also useful to get a high-level overview of the domain of interest.

Both our methods use sentence-level embeddings, rather than word-level representations, meaning that a sentence (in our case, a question) is mapped to a feature vector of a fixed size. The latter can be considered as a point in a highly dimensional space providing a numerical summary of the sentence’s meaning. In particular, here we leverage Sentence-BERT (Reimers and Gurevych, 2019), an adaptation of BERT using Siamese and triplet network structures to derive semantically meaningful sentence embeddings.

Similarity grouping. The identification of questions with similar meaning is achieved via clustering of their embeddings – an application of unsupervised learning.

Following the projection of questions into the embedding space, dimensionality reduction is first applied using Uniform Manifold Approximation and Projection for Dimension Reduction (UMAP) (McInnes, Healy and Melville, 2018). This step is meant to preserve the consistency of distance measures over the embedding space due to the curse of dimensionality (Aggarwal, Hinneburg and Keim, 2001).

In general Natural Language Processing (NLP) tasks, sentences are encoded as numerical vectors to accurately

measure the similarity among them in a distributional space. The size of these vectors, known as the dimension, determines a sort of resolution of the similarity measure. However, an excessively high dimension of embedding vectors can negate the similarity measure due to the curse of dimensionality. Thus, many Machine Learning (ML) processes typically perform a dimensionality reduction step after vector creation to mitigate the curse of dimensionality, which could adversely affect the ability of ML algorithms to distinguish among samples.

In addition, compared to other dimensionality reduction approaches, UMAP is effective for both visualization and is more effective at preserving the global structure of the data (e.g. t-SNE is often suitable for visualisation only).

Similar questions are then detected using Hierarchical Density-Based Spatial Clustering of Applications with Noise (HDBSCAN) (McInnes, Healy and Astels, 2017) on the dimensionally-reduced representations. The choice of this algorithm is motivated by the fewer hyper-parameters, its robustness to noise, and the ability to deal with variable-density clusters (e.g. competency questions forming groups of heterogeneous size). Most importantly, as a density-based method, it does not require specifying the number of clusters a priori, and does not make any assumption on the shape of clusters. This is particularly suitable for finding groups of competency questions that are automatically extracted in the previous steps, as no assumptions can be made on the number and shape of their clusters. Notably, the same approach was adopted in IDEA (de Berardinis, Carriero, Jain, Lazzari, Meroño-Peñuela, Poltronieri and Presutti, 2023) to cluster manually curated competency questions, and support participatory workflows for their refinement. The example of competency question clustering is described in Figure 7. This figure depicts the clustering results of competency questions under the 'WrittenWork' theme using UMAP. The gray cluster encompasses simpler competency questions related to object or subject identification in triples, while the yellow cluster is related to querying numerical values of birthdays for specific characters. Examination of the gray clusters reveals structural redundancy in competency questions. For example, the questions "What is the hairstyle of a fictional character?" and "What is the pseudonym of a manga character?" within the gray cluster exhibit a similar structure. Through similarity grouping based on clustering, redundancy in competency questions with similar structures can be reduced.

Paraphrase detection. Alternatively, if the goal is to filter out only those questions that have the same meaning – thereby avoiding duplicated entries, we look at paraphrase detection. One approach is to define a threshold on the cosine similarity of two sentence embeddings to deem them as semantically equivalent (e.g. considering two sentences as equivalent if they exceed a threshold t_{sim}). However, as demonstrated in Figure 7 for $t_{sim} = 0.8$, this approach is not reliable for paraphrases, as it introduces a large number of false positive and negatives.




Figure 6: Example of competency question clustering with annotations linking questions to their respective clusters for the 'WrittenWork' theme.




Figure 7: The Distribution of cosine similarities per positive (blue) and negative (orange) question pairs. The vertical blue line denotes a similarity threshold (0.8) commonly used to consider 2 sentences/embeddings as highly related.




Figure 8: Architecture of the ANN for paraphrase detection. Dotted lines denote parameter-sharing, the plus symbol stands for additive feature fusion, and the rounded box repeats the same stack of layers for 3 times before binary classification.

We address the detection of equivalent questions via transfer learning. In this case, given two questions, their sentence-level embeddings are fed to a feed-forward neural network for *paraphrase detection* – a binary classification task. The artificial neural network (ANN) is thus trained to predict whether two sentences are semantically the same (positive pairs, $y = 1$) or not (negative pairs, $y = 0$).

Despite the availability of datasets for paraphrase detection, which have been extensively used for representation learning, we did not find a pre-trained model for question-based paraphrase detection. Therefore, we implemented a deep neural network that leverages sentence-level embeddings of questions to classify them as equivalent or different. The architecture is illustrated in Figure 8. First, the sentence-level (SBERT) embeddings of two questions are passed to linear projection layers with shared parameters. This is expected to fine-tune the embeddings on the question domain,

without altering their dimension. An early fusion mechanism is then used to combine the 2 projected embeddings vectors through Hadamard product (element-wise product). The resulting vector, which still preserves the same SBERT embedding dimension, is then passed to 3 blocks of: fully-connected layers, followed by batch normalisation (Ioffe and Szegedy, 2015), ReLU activations, and dropout (Srivastava, Hinton, Krizhevsky, Sutskever and Salakhutdinov, 2014). In the final layer, a sigmoid layer is used for binary classification. By sampling from the learned Binomial distribution, the model can then be used to classify whether two questions (given their embeddings) are equivalent or not.

4. Evaluation

In this section, we describe the end-to-end and component evaluation of the RevOnt framework. The evaluated components are the Verbalisation Abstraction and the Question Generalisation stage. The end-to-end evaluation provides a meaningful quantification of system effectiveness. While, with component evaluations, it is possible to observe the impact that each component has on the system effectiveness (Resnik and Lin, 2010). One of the challenges of ontology learning and natural language processing is the dependency between the tools and techniques used in a pipeline, therefore such an evaluation would provide meaningful insight into the performance of the components and the approach overall.

4.1. Experimental setup

We have conducted a 2-step user-based experiment including $N = 15$ participants with an engineering background and familiar with ontology development. The participants are occupied mostly as ontology engineers, machine learning engineers, and data scientists. Their selection was primarily motivated by their familiarity with ontology development, especially the concept of competency questions and the participants' ability to formulate them as engineering requirements. The experiment consisted in manually reproducing the stages of verbalisation abstraction and question generation. In total, we created 20 forms³⁴ (one per theme), with a varying number of properties. The properties, which were manually reproduced, were generated using 7.6K unique triples from the WDV dataset Amaral (2022). The number of properties depends on the variety of triple verbalisations pertaining to a specific theme. Each theme falls within a domain encompassing the 20 most common types on Wikidata, ranging from written works, chemical compounds, and artists, to monuments, sport, and transportation (c.f. Table 2). Although all the WDV themes stem from Wikidata, these categories cover different domains and provide a reasonable level of diversity in terms of ontological requirements to evaluate our method.

³⁴The forms and the full responses can be found here <https://drive.google.com/drive/folders/1M7LCmqw4dc33U73GTauec02h3JNnXv4>

Given the nature of the tasks, no personal information was collected from the participants and minimal risk clearance was granted from the Research Ethics of King's College London with registration number MRSP-22/23-34232.

Table 2

An overview of the number of properties based on themes. The number of properties corresponds to the number of distinct properties of the triples of theme. The Task-Triple Ratio represents the proportion of properties relative to the triples within a specific theme.

Theme	No of properties	Property-Triple Ratio
Airport	27	7.06%
Artist	65	16.92%
Astronaut	57	16.23%
Athlete	53	13.76%
Building	67	17.4%
Celestial body	25	6.49%
Chemical compound	33	8.61%
City	72	18.79%
Comics character	79	21.01%
Food	64	17.39%
Mean of transportation	58	15.42%
Monument	62	16.31%
Mountain	23	5.98%
Painting	29	7.53%
Politician	56	14.54%
Sports team	49	12.79%
Street	21	5.46%
Taxon	27	7.01%
University	62	16.4%
Written work	21	5.45%
Total	950	

Firstly, we introduced the participants to the theme and the data that they need to complete the tasks. For each task, we provide the triple verbalisation, the subject, the subject description, the object, and the object description. The illustration used to describe the data is shown in Figure 9.

Triple verbalisation:	The violin is part of the string instrument family.
Subject:	The violin
Subject description:	A wooden chordophone in the violin family.
Object:	the string instrument family
Object description:	musical instruments that produce sound from vibrating strings

Figure 9: An illustration of example data including the triple verbalisation, the labels and descriptions of the subject and object of the triple. The example serves to explain the data needed to perform the tasks.

Next, we describe with examples the two tasks that they are required to do. The first task is to abstract a triple verbalisation given the information provided, shown in Figure 10. While the second task is to formulate questions based on the abstraction that they have created. The participants are asked to formulate three questions, where the answers to the questions will be 1) the subject of the abstraction, 2) the

property, and 3) the object of the abstraction. The tasks are represented in Figure 11.




Figure 10: Example illustration for task 1. In this illustration, we describe the first task that is to abstract the verbalisation using the given data. The abstraction is completed by generalizing the subject and the object of the triple, and not the property.




Figure 11: Example illustration for task 2. The example describes task two that is the generation of three questions, when the answer is provided. The answers are the subject, property, and object of the triple.

We arranged for each theme to be covered by two participants. In total, we have gathered 40 responses from the forms, containing approximately 1.9k annotations. Meanwhile, the WDV dataset has on average 380 triple verbalisations per theme. The coverage of each theme with annotations can be found in Table 2. We used the annotations from the user-based experiment to calculate the BLEU score, a metric for automatically evaluating machine-translated text. The BLEU score is a number between zero and one that measures the similarity of the machine-translated text to a set of high-quality reference translations Papineni, Roukos, Ward and Zhu (2002). The annotations serve as grams, each with a weight of 0.5. The number of reference human translations influences the results. Usually, more references result in better and more accurate scores. According to (Lavie, 2010) and (Google), scores over 0.3 generally reflect understandable translations and scores over 0.5 generally reflect high quality translations. In Table 3, is the interpretation of BLEU scores.

4.2. Results

In this subsection, we present the results of the evaluation of the RevOnt framework. In the following subsections, we interpret the results for each of the stages of the approach and summarise the results for the whole system. To interpret the distribution of the BLEU scores for the Verbalisation Abstraction and the Question Generation stage, we have created Box and Whisker Plots.

Table 3

Interpretation of BLEU scores

BLEU Score	Interpretation
< 0.10	Almost useless
0.10 - 0.19	Hard to get the gist
0.20 - 0.29	The gist is clear, but has significant grammatical errors
0.30 - 0.40	Understandable to good translations
0.40 - 0.50	High quality translations
0.50 - 0.60	Very high quality, adequate, and fluent translations
> 0.60	Quality often better than human

4.2.1. Verbalisation Abstraction evaluation

The distribution of the BLEU scores for the Verbalisation Abstraction stage is presented in Figure 12. As seen from the plot, the median of the scores is 0.41, which is interpreted as a high quality translation. The 75th percentile is 0.55 and the 25th percentile is 0.3. Meanwhile, the highest and the lowest data points are respectively 1 and 0. These results mean that 75% of the abstractions generated from RevOnt have a good to high quality.

During a manual thorough examination of the abstractions generated by RevOnt, that are included in the experiment, and those created by the participants we noticed that the main difference between them is that the abstraction generated by RevOnt is more general. In most cases, the system abstracts the subject and the object of the verbalisation into more general classes than a user. As shown in Example 4.1, the system abstracts Michael Jackson as a human, while a user abstracts it as a singer.

Example 4.1. *The difference between the abstraction generated by RevOnt and by the users*

Verbalisation: Michael Jackson is a member of the Michael Jackson discography.
 RevOnt: Human is a member of the discography.
 User annotation: A singer is a member of the discography.

4.2.2. Question Generation evaluation

As for the second component of RevOnt, Question Generation, we display the distribution of BLEU scores in Figure 13. The median is 0.3, which is on the border of being interpreted as a good translation. The 75th percentile is 0.47 and the 25th percentile is 0.21. The highest and the lowest data points are 0.81 and 0.1 respectively.

The non-satisfactory evaluation of this stage is explained in Figure 14. In this plot, we have evaluated individually the generation of each type of question. The red plot represents the generation of the questions when the answer is the subject of the abstraction, the yellow plot when the answer is the property of the triple, and the green plot when the answer is the object of the abstraction.

Visibly, the generation of the question when the answer is the property of the triple performs poorly. This result backs




Figure 12: Distribution of BLEU scores for the evaluation of the Verbalisation Abstraction module; considering all the domains/themes of the WDV dataset.




Figure 13: BLEU scores for the Question Generation stage. In the plot is shown the distribution of BLEU scores for the Question Generation stage for all the domains of the dataset including the three types of questions.

the argument made in Section 3 where we hypothesise that the questions that T5 generates in this case are generally identical to the ones that it generates when the answer is the object of the abstraction. In Example 4.2, we show the difference between the question that is generated by RevOnt and by the user when the answer is the property of the triple.

Example 4.2. *The difference between the questions generated by RevOnt and the users when the answer is the property of the triple.*

Verbalisation: Michael Jackson is a member of the Michael Jackson discography.
 Answer: discography/member of
 RevOnt: What is a human a member of?
 User annotation: Which is the relation between a singer and the discography?

Comparing the scenarios when the subject and the object of the abstraction are the answer, we observe that the latter is more accurate with a median of 0.42, which is a high quality translation. The 75th percentile is 0.61, the 25th percentile 0.28, and the highest point is 1.0 and the lowest is 0.08. 75%




Figure 14: BLEU score for the Question Generation stage individualized. In the plot is shown the distribution of BLEU scores for each type of question that is produced by the Question Generation stage for all the domains of the dataset.

of the scores for this category of questions is of good to high quality.

4.2.3. RevOnt 2-stage evaluation

The results of this evaluation provide further insights into the Question Generation stage. We expect a granular evaluation of the quality of each type of question to drive the refinement of the RevOnt framework. As mentioned earlier, the design choice to include all three types of questions is for experimentation only, and to provide feedback for future development.

Figure 17 reports a Box and Whisker plot of BLEU scores (generated against manually annotated questions) when RevOnt considers the *object* of the abstraction as the answer for Question Generation (which is the category of questions that performs the best). Overall, the BLEU scores for the system have a median of 0.4, which is considered as a good quality for language translation.




Figure 15: Distribution of BLEU scores for the RevOnt framework containing the Verbalisation Abstraction and the Question Generation stage with only one type of question (the question that is answered by the object of the triple).

4.2.4. Question filtering

The proposed model for question-based paraphrase detection was trained and evaluated on the Quora Question Pair (QPP) dataset (Wang, Hamza and Florian, 2017), which contains 400k+ annotated records of the form (q_1, q_2, y) with $y \in \{0, 1\}$ (0 meaning different questions; 1 meaning equivalent questions). Given the size of the dataset, we use a static split to create the training, validation, and test sets.

Although the QPP dataset covers *questions* in the wider sense (e.g. “What is the difference between a neutral and border state?”), we expect a model trained on this dataset to transfer to *competency questions*. This is motivated by the more general and easier formulation of a competency question, which is supposed to lead to less ambiguity when attempting to classify its equivalence to an alternative formulation. Therefore, the QPP dataset can be considered as a challenging benchmark to train and evaluate our model.

Our model was trained to minimise the binary cross entropy loss between the predicted and the ground-truth labels. The model is implemented in PyTorch 1.12 and trained on a NVIDIA T4. An early stopping strategy is used to prevent overfitting with 15 epochs of patience. All dropout layers are set with a drop probability of 0.5 and a fixed learning rate of 0.001 is used by the Adam optimiser with default hyper-parameters ($\beta_1 = 0.9$, and $\beta_2 = 0.999$). A threshold of 0.5 is used to sample from the learned Bernoulli distribution at the sigmoid layer (e.g. the model outputs $\hat{y} \leq 0.5$ as 0, meaning that the given questions are not equivalent; and $\hat{y} > 0.5$ as 1, in case of equivalence). This is only needed for measuring the model’s accuracy on the test set. Notably, we did not find class imbalance in QPP; hence, we can rely on the model’s accuracy as the main metric for evaluation.

After ten epochs, training stops with a test loss of 0.32 and accuracy above 86%, which is comparable to the classification results of Wang et al. (2017). The trained model can effectively identify redundant sentences, encompassing both identical sentences and semantically similar ones. For instance, multiple redundant competency questions may convey the same meaning with varied structures, such as “What is the genre of Comics?” and “Comics is in what genre?”. In sum, our model can streamline the set of competency questions by eliminating semantically identical sentences.

4.2.5. Quality features

5. Discussion and future work

This research demonstrated how the use of natural language processing methods enables the extraction of competency questions from knowledge graphs. The current implementation of RevOnt is publicly available on GitHub³⁵ and includes setup scripts and instructions to replicate all the steps of extraction process (c.f. Section 4). All code is released under the MIT license, whereas the dataset of competency questions extracted from WDV follows the Attribution 4.0 International (CC-BY 4.0).

³⁵<https://github.com/King-s-Knowledge-Graph-Lab/revont>




Figure 16: Frequency-based quality features comparison between WDV-CQ and RevOnt-CQ

To use RevOnt, users are required to produce a verbalisation of their knowledge graph using the same format of the WDV dataset (as shown in the example in Figure 5). This ensures that the question generation module produces reliable results as reported in our experiments. Nonetheless, given the modular architecture of RevOnt, users can also redefine or implement their own question generation model (e.g. using an alternative verbalisation format) and plug its outputs to the question filtering module. By doing so, we expect the hyper-parameters of the methods and tools outlined in Table 1 to remain consistent with our experimental setup.

Limitations. Overall, workflows based on language models are notoriously sensitive to the input data and the configuration of the hyper-parameters (Lavie, 2010). This issue is generally observed in the generation of the questions by RevOnt, where the distance between the triple and its verbalisation plays a significant role (c.f. Example 4.2). In addition, language models are still prone to issues and limitations which inevitably propagate into RevOnt’s pipeline. Below we highlight issues that are well-known at the intersection of natural language processing, ontology learning, and schema induction and discovery.

1. In curating the WDV dataset, Amaral (2022) already found that existing NER models performed poorly




Figure 17: Statistic-based quality features comparison between WDV-CQ and RevOnt-CQ

on specific themes (50% of the WDV dataset). Their performance is explained by the heterogeneous data used to train these models, which is comprehensively different compared to the dataset. These models are often trained with data from news articles, and emails/chat data, which explains the good performance with themes such as Artist, Athlete, City, etc. By training a NER model with data specific to the domain of interest, it is possible to make the approach independent from the Wikidata query service.

- Generally, the default language of the training data of a language model is English, which limits the usability of the model for multilingual data. Big knowledge graphs as Wikidata and DBpedia contain structured knowledge of entities in several distinct languages, and they are useful resources for cross-lingual artificial intelligence and natural language processing applications (Wang, Lv, Lan and Zhang, 2018). Therefore, language models trained with multilingual data could enhance the performance of the approach. An example of a multilingual language model is a release of BERT³⁶ that is pre-trained on the 104 languages with the largest Wikipedia pages. Another example

is T5-base³⁷, which covers English, French, Romanian, German (Raffel, Shazeer, Roberts, Lee, Narang, Matena, Zhou, Li, Liu et al., 2020).

- The dependency between the language models and methods used in NLP pipelines may also pose further issues. As the overall performance of the system is impacted by each component, the effect of errors is multiplicative. This is due to the dependency of each component to the preceding, propagating errors along the processing. In the worst cases, this may result in a final output that may be unsatisfactory despite the good performance of each component, individually Resnik and Lin (2010). In RevOnt, this has been mitigated by finding a configuration of the models (Table 1) that is optimal to the tasks and the data at hand. Hence, extending or replacing one or more components would necessitate a trial and error approach to adjust the parametrisation of our methods.

Overall, a way to mitigate the aforementioned issue is the addition of a manual validation phase across the various steps; where end-users can collaboratively refine the outputs of the system. This may be facilitated by interactive interfaces as well as crowd-sourcing and social media approaches to validate the intermediate results Asim et al. (2018). Notably, most of the issues arises from known limitations of NLP methods, especially when applied to novel or unconventional uses cases such as the one proposed in this article. Given the latest developments in NLP, we believe that most of these concerns can effectively be addressed by Large Language Models and Conversational AI systems such as ChatGPT³⁸, and Bard³⁹; which we expect to further improve the results reported in this article.

Future work. Based on the discussion above, and the opportunities opened up by RevOnt, we have also identified three directions that we plan to pursue as future work.

- The method that we have defined generates questions based on the abstraction of triple verbalisations. In particular, the question generation model (T5) performs better when is provided with the sentence and the answer. This limits the usability of the method with datasets where the answer is not explicit. It requires additional steps such as Part-Of-Speech tagging, parsing, and lemmatisation to be able to provide an entity as an answer. We plan to refine this aspect of the method by 1) detecting entities that might serve as answers or not providing an answer, and 2) providing paragraphs instead of sentences. Concerning the second point, the T-Rex dataset is a good fit, since it provides textual descriptions of Wikidata entries.
- Finally, we also plan to extend the question generation module to handle the automatic creation, or induction, of competency questions corresponding to more than

³⁶<https://huggingface.co/bert-base-multilingual-cased>

³⁷<https://huggingface.co/t5-base>

³⁸<https://openai.com/blog/chatgpt>

³⁹<https://bard.google.com>

a single triple pattern. This would allow users to extract more complex ontological requirements, beyond those emerging from a local triple-level approach.

6. Conclusions

This research is dedicated to investigating whether it is possible to extract competency questions from knowledge graphs (RQ1) and what are the quality features that an extracted question must show to be classified as a competency question (RQ2). To answer RQ1, we have defined RevOnt, an approach that can extract competency questions from knowledge graphs with the means of multiple natural language processing models. The approach is based on the reverse engineering of an ontology engineering process and uses as input the verbalisation of knowledge graph triples. RevOnt abstracts the verbalisation, generates three questions for each triple, and filters the questions. The end result of the approach is a set of core competency questions that represent the triples from where they were extracted, and the templates of the competency questions are mapped to SPARQL query templates. To address RQ2, we conducted a detailed literature review on the notion of competency questions, their function in the ontology engineering process, and the quality qualities that they demonstrate. This review concluded that the quality feature that a question must show to be classified as a competency question is its ability to be translated into a query. The feature is used as a criterion for the filtering phase of the RevOnt method.

Furthermore, we implement the RevOnt approach using the WDV dataset, a dataset of verbalisations from the Wikidata knowledge graphs. Based on a first evaluation, 75% of the abstractions generated by the first component of the approach, Verbalisation Abstraction, have a good to high quality. Meanwhile, the questions generated by the second component, Question generation, have a wider range of quality starting from poor to high. The type of question that received a higher score quality-wise is the question that RevOnt generates when the answer is the object of the triple. Lastly, the third component, Question Filtering, is evaluated with an accuracy of 86%, comparable to state-of-the-art classification tools.

The approach can be extended by fine-tuning text-to-text transformation models to accept paragraphs instead of sentences to provide more context. In addition, the model can be adjusted to not accept an answer, but alternatively use a topic detection algorithm to identify the topic of the data. Meanwhile, the performance of the approach can be improved by training the language models with the input data. This could result in Named Entity Recognition models performing better with the task and Text-to-Text transformation models generating questions with higher quality, even for non-English data.

CRedit authorship contribution statement

Fiorela Ciroku: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Resources,

Data Curation, Writing - Original Draft, Visualization. **Jacopo de Berardinis:** Methodology, Software, Validation, Formal analysis, Investigation, Resources, Data Curation, Writing- Original and Second Draft. **Jongmo Kim:** Software, Formal analysis, Resources, Writing- Second Draft. **Albert Meroño-Peñuela:** Conceptualization, Methodology, Writing - Review & Editing, Supervision, Project administration. **Valentina Presutti:** Conceptualization, Methodology, Writing - Review & Editing, Supervision, Project administration. **Elena Simperl:** Conceptualization, Methodology, Writing - Review & Editing, Supervision, Project administration.

References

- Abdelghany, A.S., Darwish, N.R., Hefni, H.A., 2019. An agile methodology for ontology development. *International Journal of Intelligent Engineering and Systems* 12, 170–181. doi:<http://dx.doi.org/10.22266/ijies2019.0430.17>.
- Aggarwal, C.C., Hinneburg, A., Keim, D.A., 2001. On the surprising behavior of distance metrics in high dimensional space, in: *Database Theory — ICDT 2001*, Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 420–434. doi:http://dx.doi.org/10.1007/3-540-44503-X_27.
- Amaral, G., 2022. WDV URL: <https://figshare.com/articles/dataset/WDV/17159045>, doi:10.6084/m9.figshare.17159045.v1.
- Amaral, G., Rodrigues, O., Simperl, E., 2022. Wdv: A broad data verbalisation dataset built from wikidata, in: *The Semantic Web – ISWC 2022*, Springer International Publishing, Cham. pp. 556–574. doi:http://dx.doi.org/10.1007/978-3-031-19433-7_32.
- Asim, M.N., Wasim, M., Khan, M.U.G., Mahmood, W., Abbasi, H.M., 2018. A survey of ontology learning techniques and applications. *Database* 2018. URL: <https://doi.org/10.1093/database/bay101>, doi:10.1093/database/bay101.
- Asprino, L., Carriero, V.A., Presutti, V., 2021. Extraction of common conceptual components from multiple ontologies, in: *Proceedings of the 11th on Knowledge Capture Conference, Association for Computing Machinery*, New York, NY, USA. p. 185–192. URL: <https://doi.org/10.1145/3460210.3493542>, doi:10.1145/3460210.3493542.
- Auer, S., Bizer, C., Kobilarov, G., Lehmann, J., Cyganiak, R., Ives, Z., 2007. Dbpedia: A nucleus for a web of open data, in: *The semantic web*. Springer, pp. 722–735. doi:https://doi.org/10.1007/978-3-540-76298-0_52.
- de Berardinis, J., Carriero, V.A., Jain, N., Lazzari, N., Meroño-Peñuela, A., Poltronieri, A., Presutti, V., 2023. The polifonia ontology network: Building a semantic backbone for musical heritage, in: *Proceedings of the 22nd International Semantic Web Conference*.
- Bezerra, C., Freitas, F., Santana, F., 2013. Evaluating ontologies with competency questions, in: *2013 IEEE/WIC/ACM International Joint Conferences on Web Intelligence (WI) and Intelligent Agent Technologies (IAT)*, pp. 284–285. doi:10.1109/WI-IAT.2013.199.
- Bottini, T., Carriero, V.A., Carvalho, J., Cathé, P., Ciroku, F., Daga, E., Daquino, M., Davy-Rigaux, A., Guillotel-Nothmann, Gurrieri, M., van Kemenade, P., Marzi, E., Meroño Peñuelala, A., Mulholland, P., Musumeci, E., Presutti, V., Scharnhorst, A., 2021. D1.1 Roadmap and pilot requirements 1st version. Technical Report. EU Commission, The Polifonia consortium.
- Čebirić, Š., Goasdoué, F., Kondylakis, H., Kotzinos, D., Manolescu, I., Troullinou, G., Zneika, M., 2019. Summarizing semantic graphs: a survey. *The VLDB Journal* 28, 295–327. doi:<https://doi.org/10.1007/s00778-018-0528-3>.
- Chen, J., Gu, J., 2021. Adol: a novel framework for automatic domain ontology learning. *Journal of Supercomputing* 77, 152–169. doi:10.1007/s11227-020-03261-7. published: 28 March 2020.

- Corcho, O., Fernández-López, M., Gómez-Pérez, A., López-Cima, A., 2005. Building Legal Ontologies with METHONTOLOGY and WebODE. Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 142–157. URL: https://doi.org/10.1007/978-3-540-32253-5_9, doi:10.1007/978-3-540-32253-5_9.
- Dietz, J.L., 2006. What is Enterprise Ontology?. Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 7–13. URL: https://doi.org/10.1007/3-540-33149-2_2, doi:10.1007/3-540-33149-2_2.
- Fadel, F., Fox, M., Gruninger, M., 1994. A generic enterprise resource ontology, in: Proceedings of 3rd IEEE Workshop on Enabling Technologies: Infrastructure for Collaborative Enterprises, pp. 117–128. doi:10.1109/ENABL.1994.330496.
- Faure, D., Poibeau, T., 2000. First experiments of using semantic knowledge learned by asium for information extraction task using intex, in: Ontology learning ECAI-2000 workshop, Citeseer. pp. 7–12.
- Fellbaum, C., 2010. WordNet. Springer Netherlands, Dordrecht. pp. 231–243. URL: https://doi.org/10.1007/978-90-481-8847-5_10, doi:10.1007/978-90-481-8847-5_10.
- Fernández-Izquierdo, A., García-Castro, R., 2022. Ontology verification testing using lexico-syntactic patterns. Information Sciences 582, 89–113. doi:<https://doi.org/10.1016/j.ins.2021.09.011>.
- Fox, M.S., Barbuceanu, M., Gruninger, M., 1996. An organisation ontology for enterprise modeling: Preliminary concepts for linking structure and behaviour. Computers in Industry 29, 123–134. URL: <https://www.sciencedirect.com/science/article/pii/0166361595000798>, doi:[https://doi.org/10.1016/0166-3615\(95\)00079-8](https://doi.org/10.1016/0166-3615(95)00079-8). wET ICE '95.
- Frantzi, K., Ananiadou, S., Mima, H., 2000. Automatic recognition of multi-word terms: the c-value/nc-value method. International journal on digital libraries 3, 115–130.
- Gangemi, A., Presutti, V., 2009. Ontology Design Patterns. Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 221–243. URL: https://doi.org/10.1007/978-3-540-92673-3_10, doi:10.1007/978-3-540-92673-3_10.
- Gick, M.L., Holyoak, K.J., 1983. Schema induction and analogical transfer. Cognitive psychology 15, 1–38. doi:[https://doi.org/10.1016/0010-0285\(83\)90002-6](https://doi.org/10.1016/0010-0285(83)90002-6).
- Google, . Evaluating models - automl translation documentation. URL: [https://cloud.google.com/translate/automl/docs/evaluate#:~:text=BLEU%20\(Bilingual%20Evaluation%20Understudy\)%20is,of%20high%20quality%20reference%20translations.](https://cloud.google.com/translate/automl/docs/evaluate#:~:text=BLEU%20(Bilingual%20Evaluation%20Understudy)%20is,of%20high%20quality%20reference%20translations.)
- Gruber, T.R., 1992. Ontolingua: A mechanism to support portable ontologies.
- Grüniger, M., Fox, M.S., 1995. The Role of Competency Questions in Enterprise Engineering. Springer US, Boston, MA. pp. 22–31. URL: https://doi.org/10.1007/978-0-387-34847-6_3, doi:10.1007/978-0-387-34847-6_3.
- Hahn, U., Romacker, M., 2001. The syndicate text knowledge base generator, in: Proceedings of the first international conference on Human language technology research.
- Ioffe, S., Szegedy, C., 2015. Batch normalization: Accelerating deep network training by reducing internal covariate shift, in: International conference on machine learning, PMLR. pp. 448–456.
- Ji, Q., Qi, G., Gao, H., Wu, T., 2019. Survey on schema induction from knowledge graphs, in: Knowledge Graph and Semantic Computing. Knowledge Computing and Language Understanding, Springer Singapore, Singapore. pp. 136–142. doi:http://dx.doi.org/10.1007/978-981-13-3146-6_12.
- John, S., Shah, N., Stewart, C., 2018. Towards a software centric approach for ontology development: Novel methodology and its application, in: 2018 IEEE 15th International Conference on e-Business Engineering (ICEBE), pp. 139–146. doi:10.1109/ICEBE.2018.00030.
- Jones, D., Bench-Capon, T., Visser, P., 1998. Methodologies for ontology development.
- Kellou-Menouer, K., Kardoulakis, N., Troullinou, G., Kedad, Z., Plexousakis, D., Kondylakis, H., 2022. A survey on semantic schema discovery. The VLDB Journal 31, 675–710. doi:10.1007/s00778-021-00717-x.
- Kouagou, N.J., Heindorf, S., Demir, C., Ngonga Ngomo, A.C., 2022. Learning concept lengths accelerates concept learning in alc, in: European Semantic Web Conference, Springer Nature Switzerland. pp. 236–252.
- Lavie, A., 2010. Evaluating the output of machine translation systems, in: Proceedings of the 9th Conference of the Association for Machine Translation in the Americas: Tutorials.
- Lenat, D., Guha, R., 1993. Building large knowledge-based systems: Representation and inference in the cyc project. Artificial Intelligence 61, 4152.
- Loper, E., Bird, S., 2002. Nltk: The natural language toolkit, in: Proceedings of the ACL-02 Workshop on Effective Tools and Methodologies for Teaching Natural Language Processing and Computational Linguistics - Volume 1, Association for Computational Linguistics, USA. p. 63–70. URL: <https://doi.org/10.3115/1118108.1118117>, doi:10.3115/1118108.1118117.
- McInnes, L., Healy, J., Astels, S., 2017. hdbscan: Hierarchical density based clustering. Journal of Open Source Software 2, 205. URL: <https://doi.org/10.21105/joss.00205>, doi:10.21105/joss.00205.
- McInnes, L., Healy, J., Melville, J., 2018. Umap: Uniform manifold approximation and projection for dimension reduction. arXiv preprint arXiv:1802.03426.
- Mintz, M., Bills, S., Snow, R., Jurafsky, D., 2009. Distant supervision for relation extraction without labeled data, in: Proceedings of the Joint Conference of the 47th Annual Meeting of the ACL and the 4th International Joint Conference on Natural Language Processing of the AFNLP: Volume 2 - Volume 2, Association for Computational Linguistics, USA. p. 1003–1011.
- Napoles, C., Sakaguchi, K., Tetreault, J., 2017. JFLEG: A fluency corpus and benchmark for grammatical error correction, in: Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers, Association for Computational Linguistics, Valencia, Spain. pp. 229–234. URL: <https://aclanthology.org/E17-2037>, doi:<http://dx.doi.org/10.18653/v1/E17-2037>.
- Nuzzolese, A.G., Gangemi, A., Presutti, V., Ciancarini, P., 2011. Encyclopedic knowledge patterns from wikipedia links, in: Aroyo, L., Welyt, C., Alani, H., Taylor, J., Bernstein, A., Kagal, L., Noy, N., Blomqvist, E. (Eds.), The Semantic Web – ISWC 2011, Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 520–536. doi:http://dx.doi.org/10.1007/978-3-642-25073-6_33.
- Öhrgren, A., Sandkuhl, K., 2005. Towards a methodology for ontology development in small and medium-sized enterprises, in: IADIS International Conference Applied Computing 2005, pp. 369–376.
- Papineni, K., Roukos, S., Ward, T., Zhu, W.J., 2002. Bleu: A method for automatic evaluation of machine translation, in: Proceedings of the 40th Annual Meeting on Association for Computational Linguistics, Association for Computational Linguistics, USA. p. 311–318. URL: <https://doi.org/10.3115/1073083.1073135>, doi:10.3115/1073083.1073135.
- Paschke, A., Schäfermeier, R., 2018. OntoMaven - Maven-Based Ontology Development and Management of Distributed Ontology Repositories. Springer International Publishing, Cham. pp. 251–273. URL: https://doi.org/10.1007/978-3-319-64161-4_12, doi:10.1007/978-3-319-64161-4_12.
- Peroni, S., 2017. A simplified agile methodology for ontology development, in: Dragoni, M., Poveda-Villalón, M., Jimenez-Ruiz, E. (Eds.), OWL: Experiences and Directions – Reasoner Evaluation, Springer International Publishing, Cham. pp. 55–69. doi:http://dx.doi.org/10.1007/978-3-319-54627-8_5.
- Piscopo, A., Simperl, E., 2018. Who models the world? collaborative ontology creation and user roles in wikidata. Proceedings of the ACM on Human-Computer Interaction 2, 1–18.
- Pouriyeh, S., Allahyari, M., Kochut, K., Arabnia, H.R., 2018. A comprehensive survey of ontology summarization: measures and methods. arXiv preprint arXiv:1801.01937.
- Presutti, V., Daga, E., Gangemi, A., Blomqvist, E., 2009. extreme design with content ontology design patterns, in: Proc. Workshop on Ontology Patterns, pp. 83–97.

- Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., Liu, P.J., et al., 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *J. Mach. Learn. Res.* 21, 1–67.
- Rajpurkar, P., Zhang, J., Lopyrev, K., Liang, P., 2016. SQuAD: 100,000+ questions for machine comprehension of text, in: *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, Association for Computational Linguistics, Austin, Texas. pp. 2383–2392. URL: <https://aclanthology.org/D16-1264>, doi:10.18653/v1/D16-1264.
- Reimers, N., Gurevych, I., 2019. Sentence-BERT: Sentence embeddings using Siamese BERT-networks, in: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, Association for Computational Linguistics, Hong Kong, China. pp. 3982–3992. URL: <https://aclanthology.org/D19-1410>, doi:10.18653/v1/D19-1410.
- Ren, Y., Parvizi, A., Mellish, C., Pan, J.Z., van Deemter, K., Stevens, R., 2014. Towards competency question-driven ontology authoring, in: *The Semantic Web: Trends and Challenges*, Springer International Publishing, Cham. pp. 752–767. doi:http://dx.doi.org/10.1007/978-3-319-07443-6_50.
- Resnik, P., Lin, J., 2010. 11 evaluation of nlp systems. *The handbook of computational linguistics and natural language processing* 57.
- Roberts, A., Raffel, C., 2020. Exploring transfer learning with t5: the text-to-text transfer transformer. Accessed on , 23–07.
- Romero, M., 2021. T5 (base) fine-tuned on squad for qg via ap. <https://huggingface.co/mrm8488/t5-base-finetuned-question-generation-ap>.
- Schekotihin, K., Rodler, P., Schmid, W., Horridge, M., Tudorache, T., 2018. Test-driven ontology development in protégé., in: *ICBO*.
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., Salakhutdinov, R., 2014. Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research* 15, 1929–1958.
- Sure, Y., Erdmann, M., Angele, J., Staab, S., Studer, R., Wenke, D., 2002. Ontoedit: Collaborative ontology development for the semantic web, in: *The Semantic Web — ISWC 2002*, Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 221–235. doi:http://dx.doi.org/10.1007/3-540-48005-6_18.
- Tettamanzi, A.G.B., Faron-Zucker, C., Gandon, F., 2014. Testing owl axioms against rdf facts: A possibilistic approach, in: Janowicz, K., Schlobach, S., Lambrix, P., Hyvönen, E. (Eds.), *Knowledge Engineering and Knowledge Management*, Springer International Publishing, Cham. pp. 519–530.
- Toutanova, K., Chen, D., 2015. Observed versus latent features for knowledge base and text inference, in: *Proceedings of the 3rd Workshop on Continuous Vector Space Models and their Compositionality*, Association for Computational Linguistics, Beijing, China. pp. 57–66. URL: <https://aclanthology.org/W15-4007>, doi:10.18653/v1/W15-4007.
- Uschold, M., Gruninger, M., 1996. Ontologies: Principles, methods and applications. *The knowledge engineering review* 11, 93–136. doi:<http://dx.doi.org/10.1017/S0269888900007797>.
- Uschold, M., King, M., 1995. Towards a methodology for building ontologies. *Citeseer*.
- Vrandečić, D., 2009. *Ontology Evaluation*. Springer Berlin Heidelberg, Berlin, Heidelberg. pp. 293–313. URL: https://doi.org/10.1007/978-3-540-92673-3_13, doi:10.1007/978-3-540-92673-3_13.
- Vrandečić, D., Krötzsch, M., 2014. Wikidata: A free collaborative knowledgebase. *Commun. ACM* 57, 78–85. URL: <https://doi.org/10.1145/2629489>, doi:10.1145/2629489.
- Wang, Z., Hamza, W., Florian, R., 2017. Bilateral multi-perspective matching for natural language sentences, in: *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence, IJCAI-17*, pp. 4144–4150. URL: <https://doi.org/10.24963/ijcai.2017/579>, doi:10.24963/ijcai.2017/579.
- Wang, Z., Lv, Q., Lan, X., Zhang, Y., 2018. Cross-lingual knowledge graph alignment via graph convolutional networks, in: *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, Association for Computational Linguistics, Brussels, Belgium. pp. 349–357. URL: <https://aclanthology.org/D18-1032>, doi:10.18653/v1/D18-1032.
- Wiśniewski, D., Potoniec, J., Ławrynowicz, A., 2021. Bigcq: A large-scale synthetic dataset of competency question patterns formalized into sparql-owl query templates. *arXiv preprint arXiv:2105.09574*.
- Wisniewski, D., Potoniec, J., Lawrynowicz, A., Keet, C.M., 2018. Competency questions and sparql-owl queries dataset and analysis. *arXiv preprint arXiv:1811.09529*.
- Zhao, L., Ichise, R., 2012. Mid-ontology learning from linked data, in: *The Semantic Web: Joint International Semantic Technology Conference, JIST 2011, Hangzhou, China, December 4-7, 2011. Proceedings 1*, Springer. pp. 112–127.